

Stochastic Bandits

February 11, 2025

1 Concentration Inequalities

Theorem 1 (Chebyshev's Inequality). *Let X be any random variable. Then for all $t > 0$,*

$$\Pr(|X - E[X]| \geq t) \leq \frac{\text{Var}(X)}{t^2}.$$

Theorem 2 (Markov inequality). *Let X be any random variable that takes only non negative values. Then for any $c > 0$,*

$$\Pr(X \geq cE[X]) \leq 1/c.$$

Theorem 3 (Hoffding's inequality). *Suppose X_1, \dots, X_n be independent bounded random variables such that for all $i \in [n]$, we have $a_i \leq X_i \leq b_i$ for all $i \in [n]$. Let $S_n = X_1 + \dots + X_n$ and so $E[S_n] = \sum_i E[X_i]$.*

For any $t > 0$, we have

$$\Pr(|S_n - E[S_n]| \geq t) \leq 2e^{-\frac{2t^2}{\sum_i (b_i - a_i)^2}}$$

If $0 \leq X_i \leq 1$ for all $i \in [n]$ then we have,

$$\Pr(|S_n - E[S_n]| \geq t) \leq 2e^{-\frac{2t^2}{n}}$$

Let $\bar{X}_n = S_n/n = \frac{X_1 + \dots + X_n}{n}$. So $E[\bar{X}_n] = \frac{\sum_i E[X_i]}{n}$ For any $t > 0$, we have

$$\Pr(|\bar{X}_n - E[\bar{X}_n]| \geq t) \leq 2e^{-\frac{2t^2 n}{\sum_i (b_i - a_i)^2}}$$

When for all i , $0 \leq X_i \leq 1$ then

$$\Pr(|\bar{X}_n - E[\bar{X}_n]| \geq t) \leq 2e^{-2t^2 n}$$

Remark. *Chebyshev's and Hoffding's inequalities works even for r.v. that can take negative values.*

2 Weak Law of Large Numbers

Let X_1, \dots, X_n be n iid (independent and identically distributed) random variables. Let $E[X] = \mu$ and $Var(X) = v$. Let $\bar{X}_n = \frac{\sum_i X_i}{n}$ be the empirical average. Note that $E[\bar{X}_n] = \mu$.

When $n \rightarrow \infty$ then the empirical average \bar{X}_n converges to the true mean μ . We can see this from Chebyshev's inequality. As X_1, \dots, X_n are independent we have

$$Var(\bar{X}_n) = Var\left(\frac{\sum_i X_i}{n}\right) = \frac{nVar(X)}{n^2} = \frac{Var(X)}{n}.$$

So as $n \rightarrow \infty$, the variance of \bar{X}_n tends to 0. So we have

$$\Pr(|\bar{X}_n - \mu| \geq \epsilon) \leq \frac{Var(X)}{n\epsilon^2} = \frac{v}{n\epsilon^2}$$

So for any fixed $\epsilon > 0$, for sufficiently large values of n (this value will depend on ϵ), the probability $\Pr(|\bar{X}_n - \mu| \geq \epsilon)$ will approach 0.

3 Stochastic Bandits

We start with the basic model of bandits, which is independent rewards. An algorithm has K possible arms to choose from, and there are T rounds, both K and T are known in advance (we will see later that one can relax the assumption that T should be known in advance). Each arm a is associated with a reward distribution D_a (pmf/pdf) which is unknown to the algorithm. The mean reward of the distribution D_a will be most relevant for us and is denoted by μ_a , i.e., $\mu_a = E_{r \sim D_a}[r]$ ($r \sim D_a$ denotes that r is sampled/drawn from the distribution D_a).

In each round $t \in T$:

1. algorithm picks an arm $a_t \in [K]$.
2. reward r_t is sampled from the distribution D_{a_t} .
3. algorithm receives the reward r_t .

The algorithm can be randomized, i.e., in any round it can fix a probability distribution over arms and can pick an arm from this distribution.

Again, it is important to note that μ_1, \dots, μ_K (and distributions D_1, \dots, D_K) are unknown to the algorithm. The algorithm only sees the rewards (of the arms pulled by the algorithm).

We make the following assumptions:

1. rewards are bounded. For simplicity, we will assume rewards at any round will be in $[0, 1]$. So the means μ_1, \dots, μ_K all are in $[0, 1]$.
2. as we are in stochastic setting, we assume all drawn rewards are independent.

An important point to note that, regardless of the algorithm is deterministic or randomized, the arms pulled by the algorithm is a random variable. Because the algorithm's decision to pull any arm in some round will depend on past history of rewards observed by the algorithm. Since the received rewards are itself a random variable, the arms pulled are also random variable. This should become clear when we will see some algorithms.

Notations: Throughout the course, we will stick to the following notations. The set of arms will be denoted by $[K]$ and the total number of rounds will be denoted by T . The best arm is the arm a that has highest mean reward. We will use a^* for the best arm and μ^* for the mean reward of the best arm. That is $a^* = \arg \max_a \mu_a$ and $\mu^* = \mu_{a^*} = \max_a \mu_a$. For any suboptimal arm, we will use $\Delta_a := \mu^* - \mu_a$ for the gap of arm a . Finally, the arm pulled by the algorithm in t th round will be denoted by a_t .

Regret: How do we measure the performance of an algorithm? Recall that the goal of the algorithm is to maximize the sum of rewards received in all rounds.

One standard approach is to compare the algorithm performance with the best possible algorithm that knows the distributions D_a for all $a \in [K]$. Note that if means μ_1, \dots, μ_K are known then best strategy to maximize total expected rewards is to always pick the arm a^* to get total expected reward of $T\mu^*$. It makes sense to define the regret of an algorithm after T rounds as

$$R(T) = \mu^*T - \sum_{t=1}^T \mu_{a_t} = \sum_{t=1}^T (\mu^* - \mu_{a_t})$$

where a_t is the arm pulled by the algorithm in the t th round. Thus $R(T)$ is the regret incurred by the algorithm after T rounds of not knowing the means $\mu_1, \mu_2, \dots, \mu_K$. Note that the regret $R(T)$ is a random variable. We are interested in the expected regret $E[R(T)]$ of the algorithm.

$$E[R(T)] = T\mu^* - E\left[\sum_{t=1}^T \mu_{a_t}\right]$$

The expectation in the above definition is taken over all the randomness, i.e., randomness over the draw of rewards from the distributions D_1, \dots, D_K and internal randomness of the algorithm (if the algorithm is randomized).

Remark. 1. By the definition only, $R(T)$ and $E[R(T)]$ of an algorithm depends on the problem-instance, i.e., means μ_1, \dots, μ_K . Of course, we want an algorithm whose expected regret is small for all instances. In the previous line, 'for all problem instances' is important. For example, consider a stupid algorithm that always pulls arm 1. This algorithm will have regret 0 if the arm 1 has mean 1 and other arms have mean 0. But of course this algorithm will suffer for other instance such as if $\mu_1 = 0$ and $\mu_2 = 1$

and will have regret of T . So we want an algorithm that performs well on all instances.

2. If the regret is small then the algorithm's performance is close to best performance when the distributions are known. So we kind of learn the distributions.
3. If r_1, \dots, r_T are rewards received by the algorithm in rounds $1, 2, \dots, T$ respectively then one can see that $E[\sum_{i=1}^T r_i] = E[\sum_{i=1}^T \mu_{a_i}]$. Hence the expected regret of an algorithm is the difference of expected total rewards of best strategy that knows μ_1, \dots, μ_K and the expected total rewards of the algorithm.

Note: In some text books, $R(T)$ is called random-regret or realized regret and $E[R(T)]$ is called Regret. Its just a matter of convention, what names should we use. Often, from the context it will be clear whether we are talking about $R[T]$ or $E[R(T)]$.

Our goal is to design an algorithm whose expected regret $E[R(T)]$ is as small as possible for all problem-instances. Like a general discussion on algorithms, we will ignore constants and only consider Big- O dependence on T . Also, we assume $T \geq K$ as any reasonable algorithm will pull each arm atleast once. Note that since reward in any round is in $[0, 1]$, $R(T) \leq T$ always. We want an algorithms for which $R(T)$ grows sublinear in T , i.e., $\frac{R(T)}{T} \rightarrow 0$ as $T \rightarrow \infty$. In other words, average regret per round should be 0 when T is large (so we essentially we have learned the unknown distributions). Smaller the regret's dependence on T , faster the rate of convergence to 0 for the average regret per round.

3.1 Explore Then Commit Algorithm

This algorithm is very intuitive and perhaps we will first come up with this algorithm. If we know the means μ_1, \dots, μ_K then in each round the best algorithm will pick the arm a^* (recall notation, a^* has highest mean $\mu^* = \max_{a \in A} \{\mu_a\}$) to maximize total expected rewards. This algorithm first try each arm m times and finds the estimate $\hat{\mu}_a$ of the mean μ_a for every arm a . Thereafter, the algorithm will always pick an arm that maximizes the empirical mean $\hat{\mu}_a$ (we will call this arm a'). The value of m is set to $\frac{T^{2/3}(\log T)^{1/3}}{K^{2/3}}$ to optimize the expected regret.

We will now prove our first theorem of this course.

Theorem 4. *For any instance (i.e., for any values of unknowns μ_1, \dots, μ_K), the expected regret of the algorithm ETC is $O(T^{2/3}(\log T)^{1/3}K^{1/3})$.*

Proof. Let $\epsilon = \sqrt{\frac{5 \ln T}{m}}$. For any arm a , let Bad_a be the event that $|\hat{\mu}_a - \mu_a| \geq \epsilon$. As $\hat{\mu}_a = \frac{\sum_{i=1}^m r_{ai}}{m}$ and all rewards are in $[0, 1]$, from Hoeffding's inequality, we have

$$\Pr(Bad_a) = \Pr(|\hat{\mu}_a - \mu_a| \geq \epsilon) \leq \frac{2}{e^{2\epsilon^2 m}} = \frac{2}{T^{10}}.$$

Algorithm 1: Explore Then Commit

1. $m = \frac{T^{2/3}(\log T)^{1/3}}{K^{2/3}};$
 2. Try each arm m times. Let for arm a , $r_{a1}, r_{a2}, \dots, r_{am}$ be the rewards received;
/* Line 2 is exploration/investment phase. In this phase,
the algorithm tries each arm so even the worst arms. But
we hope that we will be able to find near-optimal arm for
remaining rounds. */
 3. For each arm a , set $\hat{\mu}_a$ as the average received reward for the arm a ,
i.e., we have $\hat{\mu}_a = \frac{\sum_{i \in [m]} r_{ia}}{m}$. Let a' be the arm that maximizes $\hat{\mu}_a$,
i.e., $a' = \arg \max_a \{\hat{\mu}_a\};$
/* We hope a' is the near-optimal arm, i.e., $\mu^* - \mu_{a'}$ is very
small. */
 4. From the round $m \cdot K + 1$, always pick the arm a' .
-

Let $Bad = \cup_a Bad_a$ be the event that for some arm a , $|\hat{\mu}_a - \mu_a| \geq \epsilon$ holds.
So $Good = Bad^c$ is the event that, for all arms a , we have $|\hat{\mu}_a - \mu_a| < \epsilon$.
By union bound, we have

$$\Pr(Bad) \leq \sum_a \Pr(Bad_a) \leq K \cdot \frac{2}{T^{10}} \leq \frac{2}{T^9}$$

(as we are assuming $T \geq K$ as any reasonable algorithm will try each arm at least once).

Also,

$$\Pr(Good) = 1 - \Pr(Bad) \geq 1 - \frac{2}{T^9}$$

As the $\Pr(Bad)$ is negligible (we have chosen $\epsilon = \sqrt{\frac{5 \ln T}{m}}$ to ensure that this happens), it will suffice for us to only bound the expected regret conditioned on event Good. Following calculation formally show this.

$$\begin{aligned} E[R(T)] &= \Pr(Bad)E[R(T)|Bad] + \Pr(Good)E[R(T)|Good] \\ &\leq \frac{2}{T^9} \cdot T + 1 \cdot E[R(T)|Good] \\ &\leq \frac{2}{T^8} + E[R(T)|Good] \end{aligned}$$

the second inequality comes from $E[R(T)|Bad] \leq T$ as all rewards are in $[0, 1]$.

Bounding $E[R(T)|Good]$. We will now assume event $Good$, i.e., for all arms a , we have $|\hat{\mu}_a - \mu_a| < \epsilon$. Recall that $R(T) = \sum_{t=1}^T \mu^* - \mu_{a_t}$.

The contribution to regret in the investment phase can be at most mK (assuming worst case contribution of 1 in each round). The contribution to the

regret from $(mK + 1)$ st round till end is $(T - mK)(\mu^* - \mu_{a'})$ (recall that a' is the arm that maximizes $\hat{\mu}_a$ and is always chosen from $(mK + 1)$ st round) . We now claim that $(\mu^* - \mu_{a'}) < 2\epsilon$. For now let us assume this claim and bound the regret. We will prove the claim later.

$$\begin{aligned} R(T)|Good &\leq mK + (T - mK)2\epsilon \\ &\leq mK + 2T\epsilon \\ &= mK + 2T\sqrt{\frac{5 \ln T}{m}} \end{aligned}$$

As m increases mK increases and $2T\sqrt{\frac{5 \ln T}{m}}$ decreases so the above quantity will be maximized when $mK = 2T\sqrt{\frac{5 \ln T}{m}}$, i.e., when $m = (2\sqrt{5})^{2/3} \frac{T^{2/3}(\log T)^{1/3}}{K^{2/3}}$. Substituting this value of m we get $R(T)$ (conditioned on $Good$) equal to $O(T^{2/3}K^{1/3}(\log T)^{1/3})$. Obviously, we also have then $E[R(T)|Good] = O(T^{2/3}K^{1/3}(\log T)^{1/3})$.

From our earlier calculation, $E[R(T)] = \frac{2}{T^8} + E[R(T)|Good] = O(T^{2/3}K^{1/3}(\log T)^{1/3})$. Now all it remains to prove the claim that $(\mu^* - \mu_{a'}) < 2\epsilon$. Note that $\hat{\mu}_{a'} \geq \hat{\mu}_{a^*} > \mu^* - \epsilon$ and also $\mu_{a'} > \hat{\mu}_{a'} - \epsilon$. Thus we have $\mu^* - \mu_{a'} < 2\epsilon$. □

4 Successive Elimination Algorithm

One drawback of ETC algorithm is that it will continue to explore an arm large number of times (m times) even if an arm's reward history might suggest to not pull this arm further. In Successive Elimination algorithm, we discontinue the arm forever once we have belief that the arm is not good. Below is the high level description of this algorithm.

Algorithm 2: Successive Elimination - High Level Description

- 1) Pull every arm once;
 - 2) If there is 'sufficient evidence' that some arm a is not a good arm then remove this arm;
- Repeat the above steps over the remaining arms
-

Now to describe Successive Elimination fully, we just need to specify what 'sufficient evidence' is. For the same, we introduce some notations. For any arm a and round t , let $n_a(t)$ be the number of times the arm a is pulled till the round t . Obviously, we have $\sum_a n_a(t) = t$. Further, let $\hat{\mu}_a(t)$ be the empirical mean of received rewards from the arm a till round t . Formally, let r_{ai} be the reward received from the arm a on the i th pull. Then we have $\hat{\mu}_a(t) = \frac{\sum_{i=1}^{n_a(t)} r_{ai}}{n_a(t)}$. Let $\epsilon_a(t) = \sqrt{\frac{5 \log T}{n_a(t)}}$. Finally, let $UCB_a(t) = \hat{\mu}_a(t) + \epsilon_a(t)$ and $LCB_a(t) = \hat{\mu}_a(t) - \epsilon_a(t)$.

Now we describe the 'sufficient evidence' which Successive Elimination employs. Recall the analysis of ETC algorithm. There we define an event *Good* (and show that it holds with high probability) and show that conditioned on *Good*, $\hat{\mu}_a - \epsilon \leq \mu \leq \hat{\mu}_a + \epsilon$ where $\epsilon = \sqrt{\frac{5 \log T}{m}}$. Here also, we will define an event *Good* (and show it will hold with high probability) conditioned on which for all arm a and round t , we will have $LCB_a(t) \leq \mu \leq UCB_a(t)$. Now if at any time t , we have $UCB_a(t) < LCB_{a'}(t)$ for some arms a and a' then we know that $\mu_a < \mu_{a'}$. Hence, it is not a good strategy to pull arm a in any subsequent rounds because we will be better off pulling arm a' . In other words, we can eliminate the arm a for future.

Algorithm 3: Successive Elimination

```

Activate all the arms;
while #-rounds <  $T$  do
    Pull each active arm once (and receive rewards);
    Deactivate all arms  $a$  such that there exists some another arm  $a'$ 
    with  $UCB_a < LCB_{a'}$ ;
end

```

Theorem 5. For all instances, i.e., for all values of μ_1, \dots, μ_K ,

$$E[R(T)] = O(\sqrt{KT \log T})$$

Proof. Let *Good* be the following event: for all arms a and for all rounds t , we have $|\hat{\mu}_a(t) - \mu_a| < \epsilon_a(t)$ (that is, $LCB_a(t) \leq \mu_a \leq UCB_a(t)$). We will prove that

$$\Pr(\text{Good}) \geq 1 - O\left(\frac{1}{T^8}\right)$$

For now, let us assume the above. Like ETC, it will suffice to bound $E[R(T)|\text{Good}]$.

$$\begin{aligned} E[R(T)] &= \Pr(\text{Bad})E[R(T)|\text{Bad}] + \Pr(\text{Good})E[R(T)|\text{Good}] \\ &\leq O\left(\frac{1}{T^8}\right) \cdot T + 1 \cdot E[R(T)|\text{Good}] \\ &\leq O\left(\frac{1}{T^7}\right) + E[R(T)|\text{Good}] \end{aligned}$$

the second inequality comes from $E[R(T)|\text{Bad}] \leq T$ as all rewards are in $[0, 1]$.

Bounding $E[R(T)|\text{Good}]$ Each calculation in this paragraph is conditioned on *Good*. Note that $R(T) = \sum_t (\mu^* - \mu_{a_t}) = \sum_a n_a(T)(\mu^* - \mu_a)$. If we show for each arm a , $(\mu^* - \mu_a) \leq (8\sqrt{\frac{5 \log T}{n_a(T)}})$ then we are done. This is because $R(T) = \sum_a n_a(T)(\mu^* - \mu_a) \leq \sum_a n_a(T)(8\sqrt{\frac{5 \log T}{n_a(T)}}) \leq \sum_a 8\sqrt{5n_a(T) \log T}$. As \sqrt{x} is a concave function so we have $\frac{\sum_a \sqrt{n_a(T)}}{K} \leq \sqrt{\frac{\sum_a n_a(T)}{K}} = \sqrt{T/K}$. Thus $R(T) \leq 8\sqrt{5KT \log T}$. So it remains to show that for any arm a , we have $(\mu^* - \mu_a) \leq 8\sqrt{\frac{5 \log T}{n_a(T)}}$. For the sake of analysis, we will refer to the i th iteration of while loop as phase i (for any i). Our first easy observation is that the arm a^* with will never be deactivated. Let t be the last round (corresponding to the end of some phase) where the arm a remained active (in other words, arm a was played exactly once after t). Note that $n_a(t) = n_{a^*}(t)$ (because both arms a and a^* are active till t so both of them are played equal number of times, which is the number of phases completed till t). As the arm a is not deactivated at t , we must have $UCB_a(t) \geq LCB_{a^*}(t)$. Further, we have $\epsilon_a(t) = \epsilon_{a^*}(t) = \sqrt{\frac{5 \log T}{n_a(t)}} = \sqrt{\frac{5 \log T}{n_a(T)-1}}$. It is easy to now see that we have $\mu^* - \mu_a \leq 4\epsilon_a(t) = 4\sqrt{\frac{5 \log T}{n_a(T)-1}} \leq 8\sqrt{\frac{5 \log T}{n_a(T)}}$.

Bounding $\Pr(\text{Good})$ It remains to show :

$$\Pr(\text{Good}) \geq 1 - O\left(\frac{1}{T^8}\right)$$

Let $r_{a1}, r_{a2}, \dots, r_{aT}$ be T samples from the distribution D_a . Let $\text{Bad}_a(t)$ be the event that $|\frac{r_{a1} + \dots + r_{at}}{t} - \mu_a| \geq \sqrt{\frac{5 \log T}{t}}$. By Hoeffding's inequality, we have

$\Pr(Bad_a(t)) \leq O(1/e^{10 \log T}) = O(\frac{1}{T^{10}})$. Let Bad_a be the event that for some $1 \leq t \leq T$, we have $|\frac{r_{a1} + \dots + r_{at}}{t} - \mu_a| \geq \sqrt{\frac{5 \log T}{t}}$. By union bound, $\Pr(Bad_a) \leq T \cdot O(\frac{1}{T^{10}}) = O(\frac{1}{T^9})$. Let $Bad = \cup_a Bad_a$. Again by union bound, $\Pr(Bad) \leq K \cdot O(\frac{1}{T^9}) = O(\frac{1}{T^8})$. \square

The regret in the above theorem is worst-case regret, i.e, for any problem-instance (that is, for any values of K, T and μ_1, \dots, μ_K , the expected regret $E[R(T)] \leq O(\sqrt{KT \log T})$. Now we will show another type of bounds on expected regret, which will be instance -dependent bound.

Theorem 6. *The expected regret of Successive Elimination satisfies*

$$E[R(T)] \leq O(\log T) \sum_{a: \Delta_a > 0} \frac{1}{\Delta_a}$$

where $\Delta_a = \mu^* - \mu_a$ is the gap of arm a .

Proof. We define the events Bad and $Good$ as in the above theorem. Again it will suffice to bound $E[R(T)|Good]$. All calculations now are conditioned on $Good$. We claim that for any suboptimal arm a , we have $n_a(T) \leq 50000 \frac{\log T}{\Delta_a^2}$. This implies that $R(T) = \sum_a n_a(T) \Delta_a \leq O(\log T) \sum_{a: \Delta_a > 0} \frac{1}{\Delta_a}$. Suppose $n_a(T) > 50000 \frac{\log T}{\Delta_a^2}$. Consider the time t when $n_a(t) = 50000 \frac{\log T}{\Delta_a^2}$. As the arm a^* is always active we also have $n_{a^*}(t) = 50000 \frac{\log T}{\Delta_a^2}$. So now we have $\epsilon_a(t) = \epsilon_{a^*}(t) = \sqrt{\frac{5 \log T}{n_a(t)}} = \Delta_a/100$. As we have assumed $Good$, we have $LCB_{a^*}(t) \geq \mu^* - \Delta_a/100$ and $UCB_a(t) \leq \mu_a + \Delta_a/100$. So we have $UCB_a(t) < LCB_{a^*}(t)$ which implies that arm a will be eliminated in round t which contradicts that $n_a(T) > 50000 \frac{\log T}{\Delta_a^2}$. \square

UCB Algorithm

Let $n_a(t), \epsilon_a(t), \hat{\mu}_a, UCB_a(t)$ be as defined before (in Successive Elimination algorithm).

The idea of UCB is to add bonus to the empirical mean and then pick an arm that has highest value of empirical mean plus bonus. The bonus is chosen to be $\epsilon_a(t)$ which is equal to $\sqrt{\frac{5 \log T}{n_a(t)}}$ (since $\hat{\mu}_a(t) + \epsilon_a(t) = UCB_a(t)$, the algorithm, at any time t , pulls an arm that has highest value of $UCB_a(t)$).

Let us now go into the intuition behind the UCB algorithm in detail. During the initial rounds, the difference between the empirical mean and the actual mean of an arm can be significant. Consequently, selecting an arm solely based on the maximum empirical mean value is not a good strategy. To address this issue, the algorithm incorporates a bonus term. If an arm a is underexplored, the bonus is large, which encourages the algorithm to explore that arm (even if its empirical mean is small at this time). One concern might be whether the bonus term could lead the algorithm to pull bad arms excessively. However, this scenario will not happen because the bonus term diminishes as the number of pulls for the arm increases.

Algorithm 4: UCB Algorithm

```

 $UCB_a = \infty$  for all arms  $a$ ;
/* Initialization */
while #-rounds  $< T$  do
    Pull an arm that has the highest value of  $UCB_a$ ;
    /* Recall that  $UCB_a(t) = \hat{\mu}_a(t) + \epsilon_a(t)$  */
end

```

UCB achieves the same guarantee on expected regret as that of Successive Elimination. The proof is almost same so here we do not give a complete proof.

Theorem 7. *For all instances, i.e., for all values of μ_1, \dots, μ_K ,*

$$E[R(T)] = O(\sqrt{KT \log T})$$

and

$$E[R(T)] = O(\log T) \sum_{a: \Delta_a > 0} \frac{1}{\Delta_a}$$

where $\Delta_a = \mu^* - \mu_a$ is the gap of arm a .

Proof. The proof is almost same as that of Successive Elimination. To prove $E[R(T)] = O(\sqrt{KT \log T})$, we now show that for any arm a , $\mu^* - \mu_a \leq 2\sqrt{\frac{5 \log T}{n_a(T)}}$ (assuming *Good*) (and then the rest of the proof is exactly same). Let t_a be the last round when the arm a was pulled. Thus $n_a(T) = n_a(t_a)$. Note $\mu_a \geq UCB_a(t_a) - 2\epsilon_a(t_a)$ (as we are assuming *good*), $UCB_{a^*} \geq \mu^*$ (as we are assuming

Good) and $UCB_{a^*}(t_a) \leq UCB_a(t_a)$ (as arm a was pulled in round t_a). Hence, $\mu^* - \mu_a \leq 2\sqrt{\frac{5 \log T}{n_a(t_a)}} = 2\sqrt{\frac{5 \log T}{n_a(T)}}$. Now the proof goes the same as in Successive Elimination algorithm.

The proof of instance dependent bound is also similar. We here prove that (assuming *Good*) for any suboptimal arm a , we have $n_a(T) \leq 50000 \frac{\log T}{\Delta_a^2}$ (and then the proof is exactly same). Suppose not. Consider a time t when $n_a(t) = 50000 \frac{\log T}{\Delta_a^2}$. We have $\epsilon_a(t) = \Delta_a/100$. Note that for any time $t' > t$, we have $\epsilon_a(t') \leq \epsilon_a(t)$. For any time $t' > t$, we have $UCB_a(t') \leq \mu_a + 2\epsilon_a(t') \leq \mu_a + 2\epsilon_a(t) < \mu^* \leq UCB_{a^*}(t')$. This means that for any $t' > t$ we will have $UCB_{a^*}(t') > UCB_a(t')$ which means that arm a will never be pulled after t . This contradicts that $n_a(t) > 50000 \frac{\log T}{\Delta_a^2}$. □

MOSS algorithm (UCB2)

MOSS is a variant of UCB and it has the expected regret of $O(\sqrt{KT})$ and hence it beats both Successive Elimination and UCB in theory. In the next lecture, we will see that no algorithm can have smaller expected regret than $O(\sqrt{KT})$ and hence MOSS is the best possible algorithm.

MOSS is same as UCB expect how bonus is calculated. Now the bonus is set to $\sqrt{\frac{\max(\log \frac{T}{Kn_a(t)}, 0)}{n_a(t)}}$. Let

$$I_a(t) = \hat{\mu}_a(t) + \sqrt{\frac{\max(\log \frac{T}{Kn_a(t)}, 0)}{n_a(t)}}.$$

At any time t , MOSS pulls an arm that has a maximum value of I_a . One reason MOSS has better expected regret bound than UCB is that the estimation precision when $n_a(t)$ is large is more accurate in MOSS than UCB. In other words, the bonus more quickly goes to 0 in MOSS than UCB as $n_a(t)$ increases.

Algorithm 5: MOSS Algorithm

```

 $I_a = \infty$  for all arms  $a$ ;
/* Initialization */
while #-rounds  $< T$  do
    | Pull an arm that has the highest value of  $I_a$ ;
end

```

Theorem 8. *The expected regret $E[R(T)]$ of MOSS satisfies*

$$E[R(T)] = O(\sqrt{KT})$$

Proof. We will use a trick that is frequently employed in analysis of randomized algorithms. Instead of sampling from the distribution D_a at the time when the arm a is pulled, we assume that T independent samples from every distribution D_a has already been sampled before the start of the algorithm. For each arm a , let r_{a1}, \dots, r_{aT} be the T independent samples drawn from the distribution D_a . Now when the algorithm pulls arm a , we provide sample (to the algorithm) from r_{a1}, \dots, r_{aT} . In particular, the sample provided to the algorithm for i th pull of arm a is r_{ai} .

Let us define new notations. Let for any $1 \leq x \leq T$, $\hat{\mu}_{ax} = \frac{\sum_{j=1}^x r_{aj}}{x}$ be the average of first x rewards (of T samples drawn beforehand). With respect to this new notation, note that $\hat{\mu}_a(t) = \hat{\mu}_{a n_a(t)}$. Further, for any $1 \leq x \leq T$, let $I_{ax} = \hat{\mu}_{ax} + \sqrt{\frac{\max(\log \frac{T}{Kx}, 0)}{x}}$. Again note that $I_a(t) = I_{a n_a(t)}$. We will call $I_a(t)$ as the index of arm a after time t .

Let $\delta = \max\{\mu^* - \min_{1 \leq x \leq T} I_{a^* x}, 0\}$. Note that δ is a random variable. By definition only, the index of the best arm will never be less than $\mu^* - \delta$, i.e.,

$I_{a^*}(t) \geq \mu^* - \delta$ for all t . We will prove later that $E[\delta] \leq 10\sqrt{\frac{K}{T}}$. It will be helpful to keep this fact in mind.

Let us call an arm a as *Good* if $\Delta_a \leq 5\sqrt{\frac{K}{T}}$ (note that this is different - in previous algorithms, *Good* and *Bad* were events). An arm a is called *Bad* if $\Delta_a > 5\sqrt{\frac{K}{T}}$. Now it will be clear why we have defined *Good* and *Bad* arms in this way. Recall that $R(T) = \sum_a R_a(T)$ where $R_a(T) = n_a(T)\Delta_a$.

$$\begin{aligned}
R(T) &= \sum_{a:a \text{ is Good}} R_a(T) + \sum_{a:a \text{ is Bad}} R_a(T) \\
&\leq \sum_{a:a \text{ is Good}} n_a(T)5\sqrt{\frac{K}{T}} + \sum_{a:a \text{ is Bad}} R_a(T) \\
&\leq 5\sqrt{\frac{K}{T}} \sum_{a:a \text{ is Good}} n_a(T) + \sum_{a:a \text{ is Bad}} R_a(T) \\
&\leq 5\sqrt{\frac{K}{T}} \cdot T + \sum_{a:a \text{ is Bad}} R_a(T) \\
&\leq 5\sqrt{KT} + \sum_{a:a \text{ is Bad}} R_a(T)
\end{aligned}$$

Thus it suffices to show $\sum_{a:a \text{ is Bad}} R_a(T) = O(\sqrt{KT})$. Let us introduce few more notations. For any bad arm a , we define a value k_a (which is a random variable) as follows:

$$k_a = |\{1 \leq x \leq T \mid I_{a,x} > \mu_a + \frac{\Delta_a}{2}\}|$$

We also define J which is a random subset of bad arms defined as below:

$$J = \{a \in [K] \mid a \text{ is Bad and } \Delta_a > 2\delta\}$$

A very important observation is that for any arm in J , we have $n_a(T) \leq k_a$. This is the main crux of the analysis. As directly showing bounds for $E[n_a(t)]$ is difficult but later we will be able to show bounds for $E[k_a]$ for bad arms.

Now

$$\begin{aligned}
\sum_{a:a \text{ is Bad}} R_a(T) &= \sum_{a \in J} n_a(T)\Delta_a + \sum_{a \notin J: a \text{ is Bad}} n_a(T)\Delta_a \\
&\leq \sum_{a:a \text{ is Bad}} k_a \Delta_a + 2\delta T
\end{aligned}$$

Now

$$\begin{aligned}
E\left[\sum_{a:a \text{ is Bad}} R_a(T)\right] &\leq \sum_{a:a \text{ is Bad}} E[k_a]\Delta_a + 2E[\delta]T \\
&\leq \sum_{a:a \text{ is Bad}} E[k_a]\Delta_a + 20\sqrt{KT}
\end{aligned}$$

as we earlier claimed (without proof) that $E[\delta] \leq 10\sqrt{\frac{K}{T}}$. Thus it suffices to show that $\sum_{a:a \text{ is Bad}} E[k_a] \Delta_a = O(\sqrt{KT})$. Recall that for any event A , we use 1_A for indicator r.v. that takes the value 1 if A happens and 0 otherwise.

For any bad arm a , we have

$$\begin{aligned}
E[k_a] &= E[|\{1 \leq x \leq T | I_{a,x} > \mu_a + \frac{\Delta_a}{2}\}|] \\
&= E\left[\sum_{x=1}^T 1_{I_{a,x} > \mu_a + \frac{\Delta_a}{2}}\right] \\
&= \sum_{x=1}^T E[1_{I_{a,x} > \mu_a + \frac{\Delta_a}{2}}] \\
&= \sum_{x=1}^T \Pr(I_{a,x} > \mu_a + \frac{\Delta_a}{2}) \\
&= \sum_{x=1}^T \Pr(\hat{\mu}_{a,x} + \sqrt{\frac{\max(\log \frac{T}{Kx}, 0)}{x}} > \mu_a + \frac{\Delta_a}{2}) \\
&= \sum_{x=1}^T \Pr(\hat{\mu}_{a,x} - \mu_a > \frac{\Delta_a}{2} - \sqrt{\frac{\max(\log \frac{T}{Kx}, 0)}{x}}) \\
&\leq \sum_{x=1}^{8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2}} 1 + \sum_{x=8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2}}^T \Pr(\hat{\mu}_{a,x} - \mu_a > \frac{\Delta_a}{2} - \sqrt{\frac{\max(\log \frac{T}{Kx}, 0)}{x}}) \\
&= 8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2} + \sum_{x=8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2}}^T \Pr(\hat{\mu}_{a,x} - \mu_a > \frac{\Delta_a}{2} - \sqrt{\frac{\max(\log \frac{T}{Kx}, 0)}{x}})
\end{aligned}$$

As the arm a in the above calculations is bad, we have $\Delta_a > 5\sqrt{\frac{K}{T}}$. This implies that for $x \geq 8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2}$, we have

$$\begin{aligned}
\frac{\max(\log \frac{T}{Kx}, 0)}{x} &\leq \frac{\max(\log \frac{T}{K 8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2}}, 0)}{8 \frac{\log \frac{T \Delta_a^2}{K}}{\Delta_a^2}} \\
&= \frac{\Delta_a^2}{8} \cdot \frac{\log \frac{T \Delta_a^2}{8K \log(T \Delta_a^2 / K)}}{\log(T \Delta_a^2 / K)} \\
&\leq \frac{\Delta_a^2}{8}
\end{aligned}$$

The last inequality holds because $\log(T\Delta_a^2/K) > 1$ as $\Delta_a > 5\sqrt{\frac{K}{T}}$. Therefore

$$\begin{aligned} \Pr(\hat{\mu}_{a|x} - \mu_a > \frac{\Delta_a}{2} - \sqrt{\frac{\max(\log \frac{T}{Kx}, 0)}{x}}) &\leq \Pr(\hat{\mu}_{a|x} - \mu_a > \frac{\Delta_a}{2} - \frac{\Delta_a}{2\sqrt{2}}) \\ &\leq 2\exp(-2c^2\Delta_a^2x) \end{aligned}$$

where $c = \frac{1}{2}(1 - \frac{1}{\sqrt{2}})$.

Now

$$\begin{aligned} E[k_a] &\leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a^2} + \sum_{x=\frac{8 \log \frac{T\Delta_a^2}{K}}{\Delta_a^2}}^T 2\exp(-2c^2\Delta_a^2x) \\ &\leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a^2} + \sum_{x=\frac{8 \log \frac{T\Delta_a^2}{K}}{\Delta_a^2}}^{\infty} 2\exp(-2c^2\Delta_a^2x) \\ &= 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a^2} + \frac{\exp(-2c^2\Delta_a^2x_0)}{1 - \exp(-2c^2\Delta_a^2)} \end{aligned}$$

where $x_0 = 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a^2}$ and last equality comes from the geometric series summation.

As $\exp(-2c^2\Delta_a^2x_0) < 1$ we have

$$E[k_a] \leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a^2} + \frac{1}{1 - \exp(-2c^2\Delta_a^2)}$$

And hence

$$\Delta_a E[k_a] \leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a} + \frac{\Delta_a}{1 - \exp(-2c^2\Delta_a^2)}$$

Using $1 - e^{-y} \geq y - y^2/2$ we have

$$\begin{aligned} \Delta_a E[k_a] &\leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a} + \frac{\Delta_a}{2c^2\Delta_a^2 - 2c^4\Delta_a^4} \\ &\leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a} + \frac{1}{2c^2\Delta_a(1 - c^2\Delta_a^2)} \\ &\leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a} + \frac{1}{2c^2\Delta_a(1 - c^2)} \\ &\leq 8 \frac{\log \frac{T\Delta_a^2}{K}}{\Delta_a} + \frac{1}{2c^2(1 - c^2)} \frac{\sqrt{T}}{5\sqrt{K}} \end{aligned}$$

One can check that the maximum value of $f(y) = \frac{\log(Ty^2/K)}{y}$ is $O(\sqrt{\frac{T}{K}})$.
Thus

$$\sum_{a:a \text{ is bad}} \Delta_a E[k_a] \leq \sum_{a:a \text{ is bad}} O\left(\frac{\sqrt{T}}{\sqrt{K}}\right) = O(\sqrt{KT})$$

All remains to show: $E[\delta] \leq 10\sqrt{\frac{K}{T}}$. □

5 Lower Bounds

We started with ETC and show its expected regret of $O(T^{2/3}K^{1/3}(\log T)^{1/3})$. Then we show that the expected regret of Successive Elimination and UCB is $O(\sqrt{KT \log T})$. Finally, we show that MOSS beats UCB and Successive Elimination and has expected regret of $O(\sqrt{KT})$. What next? Can we have an algorithm that has even better regret than MOSS, i.e., $O(\sqrt{KT})$? What if researchers are unable to find a better regret bound then can we conclude that better bound is not possible? Surely not. It could be lack of understanding or knowledge, preventing us from coming up with a better algorithm.

Now our focus will be lower bounds, i.e., we will show that no algorithm can have better regret bound than some threshold. In fact we will prove the following theorem

Theorem 9. *No algorithm (deterministic or randomized) can have expected regret of $o(\sqrt{KT})$ for all instances.*

Note that the possibilities of algorithms are infinite. An algorithm can perhaps do anything. Formally, a deterministic algorithm is any function that takes as input K, T and for any $1 \leq t \leq T$ maps a sequence (of pairs consisting of arm and received reward) $\{(a_1, r_1), (a_2, r_2), \dots, (a_{t-1}, r_{t-1})\}$ to some arm $a_t \in [K]$. Since rewards can take any value in $[0, 1]$, the number of distinct algorithms is infinite. Thus, it is not possible to prove the above theorem by showing that a particular algorithm does not have a better regret bound for all such algorithms (as the number of such algorithms is infinite).

Not surprisingly, we will need new tools to prove such lower bound. We now will study elementary tools from Information theory which will be sufficient to prove the above theorem. In fact we will not directly prove the Theorem 9. First we will consider a toy problem (coin testing problem of exercise) and show a lower bound for this problem. This itself will require new tools from Information Theory. Then after getting some familiarity with these tools, we will prove the Theorem 9.

Testing Coin(ϵ, δ) Problem: It is promised that a given coin is either fair, i.e., $p = 1/2$ or ϵ -biased, i.e., $p = 1/2 + \epsilon$ (algorithm knows ϵ). Give an algorithm A that takes samples from that coin (as input) and has the following guarantee:

- If $p = 1/2$, algorithm must say ‘Fair’ with at least $1 - \delta$ probability.
- If $p = 1/2 + \epsilon$, the algorithm must say ‘Biased’ with at least $1 - \delta$ probability.

From Hoeffding’s inequality, it is easy to show following:

Lemma 1. *There is a deterministic algorithm that solves Testing coin problem and needs only $O(1/\epsilon^2)$ samples for any constant $\delta < 1/2$.*

We will now show the following lower bound:

Lemma 2. *Let $\epsilon \leq 1/3$. Any deterministic algorithm will need at least $\frac{(1-2\delta)^2}{5\epsilon^2}$ samples to solve Testing Coin(ϵ, δ) problem.*

Now the goal is to prove Lemma 2. As said before, we begin with understanding new tools.

Total Variation Distance and KL divergence

Recall that a probability distribution P over a finite domain Ω is a function $P : \Omega \rightarrow [0, 1]$ that maps each $\omega \in \Omega$ to a real in $[0, 1]$ satisfying:

1. $P(\omega) \geq 0$ for all $\omega \in \Omega$
2. $\sum_{\omega \in \Omega} P(\omega) = 1$.

Now we will define the notion of distance between two probability distributions. Any notion of distance d must necessarily satisfy these two properties - $d(P, P) = 0$ for any distribution P and $d(P, Q) > 0$ for any two distinct distribution P and Q . There can be many ways to define distance. for instance, we can consider Euclidean distance or ℓ_2 norm: $d(P, Q) = \sum_{\omega} (P(\omega) - Q(\omega))^2$ or we can consider ℓ_1 norm called as total variation distance $d(P, Q) = \frac{1}{2} \sum_{\omega} |P(\omega) - Q(\omega)|$. Each notion of distance has its own merits and limitaions. For us, the notion of total variation distance also known as statistical distance will be important.

Definition 1. For any two distributions P and Q defined on a finite domain Ω , the total variation distance between P and Q , denoted as $d_{tv}(P, Q)$ is defined as

$$d_{tv}(P, Q) = \frac{1}{2} \sum_{\omega \in \Omega} |P(\omega) - Q(\omega)|$$

It is easy to see that:

1. $d_{tv}(P, Q) = 0$ if $P = Q$.
2. $d_{tv}(P, Q) > 0$ for any $P \neq Q$.
3. $d_{tv}(P, Q) = d_{tv}(Q, P)$.
4. $d_{tv}(P, Q) \leq d_{tv}(P, R) + d_{tv}(R, Q)$ (triangle inequality)

Now we are going to see one of the most useful and remarkable property of total variation distance.

Lemma 3.

$$d_{tv}(P, Q) = \max_{S \subseteq \Omega} |P(S) - Q(S)| \tag{1}$$

where note that $P(S) = \sum_{\omega \in S} P(\omega)$ and $Q(S) = \sum_{\omega \in S} Q(\omega)$.

Now we define KL divergence between two distributions.

Definition 2. The KL divergence $KL(P, Q)$ is defined as

$$KL(P, Q) = \sum_{\omega \in \Omega} P(\omega) \ln \frac{P(\omega)}{Q(\omega)}$$

The KL divergence satisfy:

Lemma 4. • $KL(P, P) = 0$

• $KL(P, Q) > 0$ if $P \neq Q$

However, note that $KL(P, Q) \neq KL(Q, P)$ and also it does not satisfy the triangle inequality. That is why it is called divergence not distance. The usefulness of KL divergence comes from the fact that the KL divergence of a product distribution just adds up. Let us see first what is a product distribution and then the utility of KL divergence (I have talked about the product distribution while discussing independent trials).

Definition 3. Given two distributions P and Q on the domains Ω_1 and Ω_2 respectively, then $P \times Q$ is a distribution on domain $\Omega_1 \times \Omega_2$ such that

$$(P \times Q)(\omega_1, \omega_2) = P(\omega_1)Q(\omega_2) \quad \forall \omega_1 \in \Omega_1, \omega_2 \in \Omega_2$$

In general, if P_i is a distribution on Ω_i then the product distribution $P_1 \times P_2 \times \dots \times P_m$ is a distribution on domain $\Omega_1 \times \dots \times \Omega_m$ such that

$$(P_1 \times P_2 \times \dots \times P_m)(\omega_1, \omega_2, \dots, \omega_m) = P_1(\omega_1)P_2(\omega_2) \dots P_m(\omega_m)$$

We will use, for any distribution P , $P^{\otimes 2}$ for $P \times P$. In general, we have $P^{\otimes m} = P \times \dots \times P$ (m -times).

The product distribution corresponds to independent trials. Drawing m independent samples from a distribution P is exactly same as drawing one sample from $P^{\otimes m}$. Now we will see why KL divergence is extremely useful.

Lemma 5. If for all $i \in [m]$, P_i and Q_i are distributions on Ω_i then

$$KL(P_1 \times P_2 \times \dots \times P_m, Q_1 \times Q_2 \times \dots \times Q_m) = KL(P_1, Q_1) + KL(P_2, Q_2) + \dots + KL(P_m, Q_m)$$

Corollary 1.

$$KL(P^{\otimes m}, Q^{\otimes m}) = m \cdot KL(P, Q)$$

The above property is not true for total variation distance. We do not have $d_{tv}(P^{\otimes m}, Q^{\otimes m}) = m d_{tv}(P, Q)$. In fact, other than $d_{tv}(P^{\otimes m}, Q^{\otimes m}) \leq m d_{tv}(P, Q)$, we do not know of any other inequality between these two.

We have now seen both total variation distance and KL divergence, each having its own merits and limitations. Now we will see that these two are connected by an inequality, which forms the basis of powerful application.

Lemma 6 (Pinsker inequality). For any two probability distributions P and Q , we have

$$d_{tv}(P, Q) \leq \sqrt{\frac{1}{2} KL(P, Q)}$$

Following is also useful.

Lemma 7 (The Bretagnolle–Huber bound).

$$d_{tv}(P, Q) \leq \sqrt{1 - e^{-KL(P, Q)}} \leq 1 - \frac{1}{2} e^{-KL(P, Q)}$$

Now we have all the necessary tool to prove Lemma 2. First we will prove the lemma for deterministic algorithms.

Proof of Lemma 2 Suppose there is a deterministic algorithm A that solves the Testing Coin problem and uses m independent coin tosses. Note that algorithm A is just a function $A : \{H, T\}^m \rightarrow \{\text{Fair}, \text{Biased}\}$ that maps a sequence of head and tails of length m to Fair or Biased. We partition $\{H, T\}^m$ into two subsets as follows: $A_f = \{y \in \{0, 1\}^m : A(y) = \text{Fair}\}$ and $A_b = \{y \in \{0, 1\}^m : A(y) = \text{Biased}\}$, i.e., A_f consists of all those sequence which gets mapped to Fair and A_b consists of all those sequences that gets mapped to Biased.

Let $p = \Pr(\text{Head})$. Let $(\frac{1}{2})^{\otimes m}$ and $(\frac{1}{2} + \epsilon)^{\otimes m}$ be the product distributions corresponding to m independent coin tosses with $p = 1/2$ and $p = 1/2 + \epsilon$ respectively. For convenience, let $D_f = (\frac{1}{2})^{\otimes m}$ and $D_b = (\frac{1}{2} + \epsilon)^{\otimes m}$. We will bound tv distance between $D_f = (\frac{1}{2})^{\otimes m}$ and $D_b = (\frac{1}{2} + \epsilon)^{\otimes m}$ in two different ways.

Note that $D_f(A_f) = \sum_{\omega \in A_f} D_f(\omega) \geq 1 - \delta$ and $D_b(A_f) = \sum_{\omega \in A_f} D_b(\omega) \leq \delta$ (because the algorithm A solves the Testing Coin(ϵ, δ) problem). Thus by Lemma 3,

$$d_{tv}(D_f, D_b) \geq D_f(A_f) - D_b(A_f) \geq 1 - 2\delta.$$

Now let us bound tv distance using Pinsker inequality. As said earlier, calculation of KL divergence between D_f and D_b is easy as they both are product distributions.

$$KL((\frac{1}{2})^{\otimes m}, (\frac{1}{2} + \epsilon)^{\otimes m}) = m \cdot KL(1/2, 1/2 + \epsilon)$$

where $KL(1/2, 1/2 + \epsilon)$ is the KL divergence between the distributions $(1/2, 1/2)$ and $(1/2 + \epsilon, 1/2 - \epsilon)$.

$$\begin{aligned} KL(1/2, 1/2 + \epsilon) &= \frac{1}{2} \ln \frac{1}{2(\frac{1}{2} + \epsilon)} + \frac{1}{2} \ln \frac{1}{2(\frac{1}{2} - \epsilon)} \\ &= \frac{1}{2} \ln \frac{1}{4(1/4 - \epsilon^2)} = \frac{1}{2} \ln \frac{1}{(1 - 4\epsilon^2)} \\ &= \frac{1}{2} (4\epsilon^2 + \frac{(4\epsilon^2)^2}{2} + \frac{(4\epsilon^2)^3}{3} + \dots) = 2\epsilon^2(1 + 2\epsilon^2 + 16\epsilon^4/3 + \dots) \leq 10\epsilon^2 \end{aligned}$$

for $\epsilon \leq 1/3$ (probably 10 can be replaced by a smaller number but as we are not optimizing constants I put here bigger number so that the inequality is easy to derive).

Now we will use Pinsker inequality to bound the tv distance.

$$\begin{aligned} d_{tv}((\frac{1}{2})^{\otimes m}, (\frac{1}{2} + \epsilon)^{\otimes m}) &\leq \sqrt{\frac{1}{2} KL((\frac{1}{2})^{\otimes m}, (\frac{1}{2} + \epsilon)^{\otimes m})} \\ &\leq \sqrt{5m\epsilon} \end{aligned}$$

Now we have

$$1 - 2\delta \leq d_{tv}(D_f, D_b) \leq \sqrt{\frac{1}{2} KL((\frac{1}{2})^{\otimes m}, (\frac{1}{2} + \epsilon)^{\otimes m})} \leq \sqrt{5m\epsilon}$$

which implies that $m \geq \frac{(1-2\delta)^2}{5\epsilon^2}$.

Lower Bound for randomized algorithms

Lemma 8. *Let $\epsilon \leq 1/3$. The expected number of coin tosses by any randomized algorithm that solves $\text{Testing Coin}(\epsilon, \delta)$ problem is at least $\frac{(1-6\delta)^2}{15\epsilon^2}$.*

In the class, I discussed that any randomized algorithm R is just a distribution over some deterministic algorithms, i.e., for any randomized algorithm R , there exists deterministic algorithms A_1, \dots, A_s (for some finite s) and a distribution (q_1, \dots, q_s) over these deterministic algorithms such that R picks a A_i ($i \in [s]$) (with probability q_i) and then run the deterministic algorithm A_i on the given input.

Proof of Lemma 8 Let R be a randomized algorithm that solves the Testing Coin (ϵ, δ) using at most $m = \frac{(1-6\delta)^2}{15\epsilon^2}$ expected number of coin tosses. Let A_1, \dots, A_s be deterministic algorithms each requiring at most m_1, m_2, \dots, m_s coin tosses and let R picks each A_i with probability q_i . Now consider a table with two rows and s columns — one row corresponding to $p = 1/2$ (called row 1) and one corresponding to $p = 1/2 + \epsilon$ (called row 2), and s columns corresponding to each deterministic algorithms A_1, \dots, A_s .

We put \checkmark in the cell corresponding to row 1 and column A_j if $\Pr(A_j \text{ returns 'Biased' when } p = 1/2) \leq 3\delta$. Otherwise we put \times , i.e., if $\Pr(A_j \text{ returns 'Biased' when } p = 1/2) > 3\delta$. For row 2, we put \checkmark in the cell corresponding to row 2 and column A_j if $\Pr(A_j \text{ returns 'Fair' when } p = 1/2 + \epsilon) \leq 3\delta$. Otherwise we put \times , i.e., if $\Pr(A_j \text{ returns 'Fair' when } p = 1/2 + \epsilon) > 3\delta$. By $\text{cell}(i, A_j)$ ($i \in \{1, 2\}, j \in [s]$), we mean the symbol (\checkmark or \times) in the cell corresponding to row i and column A_j in the above table.

Now,

$$\begin{aligned} \delta &\geq \Pr(R \text{ returns 'Biased' when } p = 1/2) \quad (\text{because } R \text{ solves Testing Coin}(\epsilon, \delta)) \\ &= \sum q_j \Pr(A_j \text{ returns 'biased' when } p = 1/2) \\ &\geq \sum_{j: \text{cell}(1, A_j) = \times} q_j 3\delta \\ &\implies \sum_{j: \text{cell}(1, A_j) = \times} q_j \leq 1/3 \end{aligned}$$

Similarly,

$$\sum_{j: \text{cell}(2, A_j) = \times} q_j \leq 1/3$$

This means that $\sum_{j: \text{cell}(1, A_j) = \checkmark \text{ and } \text{cell}(2, A_j) = \checkmark} q_j \geq 1 - (1/3 + 1/3) \geq 1/3$. That is, R puts at least $1/3$ probability on those A_j 's for which both $\Pr(A_i \text{ returns 'Biased' when } p = 1/2) \leq 3\delta$ and $\Pr(A_i \text{ returns 'Fair' when } p = 1/2 + \epsilon) \leq 3\delta$ are true. Let $J = \{j : \text{cell}(1, A_j) = \checkmark \text{ and } \text{cell}(2, A_j) = \checkmark\}$. Note that for any $j \in J$, A_j solves Testing Coin (ϵ, δ) problem and is a deterministic algorithm. By Lemma 2,

$m_j \geq \frac{(1-2(3\delta))^2}{5\epsilon^2} = \frac{(1-6\delta)^2}{5\epsilon^2}$ (recall m_j is the number of coin tosses by A_j). Therefore, the expected number of coin tosses by R , is $\sum_j q_j m_j \geq \frac{1}{3} \cdot \frac{(1-6\delta)^2}{5\epsilon^2} = \frac{(1-6\delta)^2}{15\epsilon^2}$.

Remark. I presented the Lemma 8 and its proof in the class in a slightly different way. The Lemma 8 is more stronger than the one I stated in the class. There, I assumed that the running time (or the number of samples) of R is a deterministic so each A_1, \dots, A_s uses fixed number of samples. As the above proof shows, we can show lower bound on the expected number of samples (by same proof).

Lower Bounds for Bandits

We start with showing lower bounds when there are only two arms. The best upper bound in this case is $O(\sqrt{T})$. We show a matching lower bound.

Theorem 10. No algorithm can achieve an expected regret of $O(T^{0.5-\alpha})$ for any constant $\alpha > 0$ for all instances.

Proof. Let there exists an algorithm B that has an expected regret of $O(T^{0.5-\alpha})$ for some $\alpha > 0$ for all instances. We will show that, using B , we can get a randomized algorithm R that solves Testing Coin($\epsilon, 1/50$) in $o(1/\epsilon^2)$ coin tosses contradicting Lemma 8.

Let $\epsilon < 1/3$ be the input of Testing Coin(ϵ, δ) problem. Consider a following bandit instance:

- $K = 2$ (there are two arms)
- Arm 1's reward takes value either 0 or 1. The probability of reward being 1 is the probability that the coin (of Testing Coin($\epsilon, 1/50$)) outputs a Head.
- Arm 2's reward distribution is also a Bernoulli with mean reward of $\frac{1}{2} + \frac{\epsilon}{2}$, i.e., reward takes either 0 or 1 and the probability of reward being one is $\frac{1}{2} + \frac{\epsilon}{2}$.
- $T = (200/\epsilon)^{\frac{1}{0.5+\alpha}}$

Note in the above that $T = o(1/\epsilon^2)$. Now we describe a randomized algorithm R . The algorithm R runs B on the above bandit instance (recall that we assumed there exists an algorithm B that has an expected regret of $O(T^{0.5-\alpha})$ for some $\alpha > 0$ for all bandit instances). Whenever B pulls arm 2, the algorithm R generates a sample from the arm 2's reward distribution (Bernoulli with mean reward of $\frac{1}{2} + \frac{\epsilon}{2}$) and give it to the bandit algorithm B (these are internal randomness of the randomized algorithm R). Whenever B pulls arm 1, R uses coin toss of the input coin (of Testing Coin(ϵ, δ)) to provide the reward to B — if the coin toss outputs Head, it provides a reward of 1 to B and if the result of coin toss is Tail, R provides a reward of 0 to B . Finally, R returns 'Fair' if B (after T rounds) pulls the arm 2 at least $T/2$ times and otherwise return 'Biased' (i.e., if B pulls arm 1 more than $T/2$ times).

Note that R uses at most $T = o(1/\epsilon)^2$ coin tosses. Now we prove correctness of the algorithm R — if the coin is fair, i.e., $p = \Pr(\text{Head}) = 1/2$ then B pulls the arm 2 more than $T/2$ times (and hence outputs ‘Fair’) with at least $49/50$ probability and if the coin is biased, i.e., $p = 1/2 + \epsilon$ then B pulls the arm 1 more than $T/2$ times (and hence outputs ‘Biased’) with at least $49/50$ probability. In other words, B will always pull the best arm $> T/2$ times with at least $49/50$ probability. To see this, note that in both cases $E[R(T)] = \frac{\epsilon}{2} E[N_s]$ where N_s is the number of times B pulls the suboptimal arm (which is arm 1 when $p = 1/2$ and arm 2 when $p = 1/2 + \epsilon$). Now we have

$$E[R(T)] = \frac{\epsilon}{2} E[N_s] \leq T^{\frac{1}{2}-\alpha} \implies E[N_s] \leq \frac{T^{1/2+\alpha}}{100} T^{\frac{1}{2}-\alpha} = T/100$$

By Markov’s inequality, we have $\Pr(N_s \geq T/2) \leq 1/50$. Thus, R solves Testing Coin $(\epsilon, 1/50)$ problem in $T = o(1/\epsilon)^2$ coin tosses which contradicts Lemma 8. \square

Now our goal will be to show the lower bound for K arms.

Theorem 11. *No algorithm can have an expected regret of $\leq \sqrt{KT}/5000$ on all bandit instances.*

To prove the above theorem, we first need to define Biased-Coin Identification problem.

Biased-Coin Identification: Input is K coins C_1, \dots, C_K and $\epsilon > 0$. It is promised that there is exactly one biased coin for which $\Pr(\text{Head}) = 1/2 + \epsilon$ and rest of the others are fair coin. The goal of the algorithm is to return the index of the unfair coin with probability at least $4/5$.

Formally, Let $P^i = (\frac{1}{2}, \dots, \frac{1}{2} + \epsilon, \dots, \frac{1}{2})$ be K -tuple where all coordinate is $1/2$ except the i th one which is $1/2 + \epsilon$. If the tuple corresponding to the input coins probabilities is P^i then the algorithm should return i with at least $4/5$ probability, i.e., $\Pr_{P^i}(y = i) \geq 4/5$ where y denotes the output of the algorithm and $\Pr_{P^i}(y = i)$ denotes the probability of the event $y = i$ when the input coins probability tuple is P^i .

Theorem 12. *Let $\epsilon \leq 1/3$. Any deterministic algorithm requires $> K/1000\epsilon^2$ number of total coin tosses to solve Biased-Coin Identification problem.*

We will prove the above theorem later. First we prove the Theorem 11 assuming the above theorem.

Proof of Theorem 11 (assuming Theorem 12)

Proof. Suppose there is such an algorithm B . We now show that, using B , we can solve Biased-coin Identification problem in $K/2000\epsilon^2$ total coin tosses contradicting Theorem 12.

We run B on a bandit instance on K arms and $T = K/2000\epsilon^2$. Whenever B pulls an arm j , we toss the C_j th coin and provide to B a reward of 1 if the coin

toss results in Head and a reward of 0 if the coin toss result is Tail. Finally, we return the index of the arm that has been pulled highest number of times.

It is easy to see that if i th coin is biased then the reward distributions (of the bandit instance given as input to B) of all arms except the i th one is a Bernoulli with a mean of $1/2$ and the reward distribution of the i th arm is a Bernoulli with a mean of $1/2 + \epsilon$.

Note that for any of the above bandit instances, we have $E[R(T)] = \epsilon(T - E[N_o])$ where N_o is the number of times the optimal arm is pulled by the algorithm. By our assumption, $E[R(T)] \leq \frac{\sqrt{KT}}{5000} \implies T - E[N_o] \leq \frac{\sqrt{KT}}{5000\epsilon} = \frac{T}{100}$. By Markov inequality, $\Pr(T - N_o \geq T/2) \leq 1/50$. In other words, with at least $49/50$ probability, B will pull the optimal arm strictly more than $T/2$ times (which also implies that the optimal arm will be pulled highest number of times). Hence with at least $49/50 > 4/5$ probability, the output of the algorithm will be the index of the optimal arm which is also the index of the biased coin. Moreover, our algorithm uses only $K/2000\epsilon^2$ total coin tosses. \square

Proof of Theorem 12

Let A be an algorithm that solves the Biased Coin Identification problem in $m < \frac{K}{1000\epsilon^2}$ total coin tosses.

We will now see that A can be viewed as a complete binary tree of height m where recall that m is the number of total coin tosses by the algorithm. All nodes of this tree are labelled by some index $i \in [K]$. From any internal node, the edge to left child is labelled by H and the edge to right child is labelled by T (this ordering is arbitrary). The algorithm starts from root and reach a leaf. When the algorithm reaches any internal node labelled by i then it means that the algorithm tosses the coin C_i . If the result of coin toss is H then the algorithm moves to the left child and if the result is T then it goes to the right child. In this way, algorithm starts from the root node and reaches some leaf node. If the algorithm reaches a leaf with label i then it returns i . Below we define some notions related with this tree.

Notations: The label of a node v will be denoted by $L(v)$ (which will be in $[K]$). The height of a node v will be denoted by $h(v)$. By convention we assume height of the root node is 0 and so the height of the leaf nodes are m . For any node v , we will use $\ell(v)$ for left child of v and $r(v)$ for right child of v . The set of all nodes at height h will be denoted by V_h and moreover $V_h = (v_{h1}, v_{h2}, \dots, v_{h2^h})$. Note that $|V_h| = 2^h$.

Let y denote the output of the algorithm (which will be in $[K]$). Recall that $P^i = (\frac{1}{2}, \dots, \frac{1}{2} + \epsilon, \dots, \frac{1}{2})$ is the K -tuple where all coordinate is $1/2$ except the i th one which is $1/2 + \epsilon$. For any node v , we use $\Pr_Q(v)$ for the probability that the algorithm reaches node v when the input coins probability tuple is Q . Moreover, $p_Q(v, \ell(v))$ is the probability that the algorithm reaches $\ell(v)$ (left child of v) given that it is currently at v when the input coins probability tuple is Q . Similarly, $p_Q(v, r(v))$ is the probability that the algorithm reaches $r(v)$

(right child of v) given that it is currently at v when the input coins probability tuple is Q . Let $D_Q^h = (\Pr_Q(v_{h1}), \dots, \Pr_Q(y_{h2^h}))$ be the distribution induced on the nodes at height h , when the input coins probability tuple is Q . Note that $\sum_{k=1}^{2^h} \Pr_Q(v_{hk}) = 1$ for all Q (as the algorithm will reach exactly one node at any height) so this is a probability distribution. Finally, the output of the algorithm will be denoted by y .

Since A solves the Biased Coin Identification problem, we must have $\Pr_{P^i}(y = i) \geq 4/5$ for all $i \in [K]$.

Claim 1. $d_{tv}(D_{P^i}^m, D_{P^j}^m) \geq 3/5$ for all $i \neq j$.

Proof. Consider set $S_i = \{v \in V_m : L(v) = i\}$ and $S_j = \{v \in V_m : L(v) = j\}$ (S_i is the set of all leaves with label i . Similarly S_j is the set of all leaves with label j). We have $D_{P^i}^m(S_i) = \sum_{v \in S_i} D_{P^i}^m(v) = \Pr_{P^i}(y = i) \geq 4/5$ (because $\Pr_{P^i}(y = i) = \sum_{v \in S_i} \Pr_{P^i}(v) = \sum_{v \in S_i} D_{P^i}^m(v)$). Similarly, $D_{P^j}^m(S_j) \geq 4/5$ and hence $D_{P^j}^m(S_i) \leq 1/5$. By Lemma 3, we have $d_{tv}(D_{P^i}^m, D_{P^j}^m) \geq |D_{P^i}^m(S_i) - D_{P^j}^m(S_i)| \geq 3/5$. □

Now we would like to bound $KL(D_{P^i}^m, D_{P^j}^m)$. Unfortunately, bounding this directly is not easy. We will use $F = (1/2, \dots, 1/2)$ (all $1/2$) as a reference distribution. Let us calculate $KL(D_F^m, D_{P^j}^m)$. Our claim is that

$$KL(D_F^m, D_{P^j}^m) = KL(D_F^{m-1}, D_{P^j}^{m-1}) + \sum_{v \in V_{m-1}} \frac{1}{2^{m-1}} KL((1/2, 1/2), (p_{P^j}(v, \ell(v)), p_{P^j}(v, r(v))))$$

Before proving the above let us interpret this. The KL divergence of the distributions induced on nodes at level m by F and P^j is equal to the divergence of the distributions induced on nodes at level $m-1$ plus divergence between them conditioned on the fact that algorithm is currently at height $m-1$. *For clarity, we will use $p_j(v, \ell(v))$ for $p_{P^j}(v, \ell(v))$ and $p_j(v, r(v))$ for $p_{P^j}(v, r(v))$ from now on.*

Let us prove the above.

$$\begin{aligned}
KL(D_F^m, D_{P_j}^m) &= \sum_{v \in V_m} \frac{1}{2^m} \ln \frac{1}{2^m \Pr_{P_j}(v)} \\
&= \sum_{v \in V_{m-1}} \frac{1}{2^m} \ln \frac{1}{2^m \Pr_{P_j}(\ell(v))} + \frac{1}{2^m} \ln \frac{1}{2^m \Pr_{P_j}(r(v))} \\
&= \sum_{v \in V_{m-1}} \frac{1}{2^m} \ln \frac{1}{2^m \Pr_{P_j}(v) \cdot p_j(v, \ell(v))} + \frac{1}{2^m} \ln \frac{1}{2^m \Pr_{P_j}(v) \cdot p_j(v, r(v))} \\
&= \sum_{v \in V_{m-1}} \frac{1}{2^m} \ln \frac{1}{2^{m-1} \Pr_{P_j}(v)} + \sum_{v \in V_{m-1}} \frac{1}{2^m} \ln \frac{1}{2 p_j(v, \ell(v))} \\
&\quad + \sum_{v \in V_{m-1}} \frac{1}{2^m} \ln \frac{1}{2^{m-1} \Pr_{P_j}(v)} + \sum_{v \in V_{m-1}} \frac{1}{2^m} \ln \frac{1}{2 p_j(v, r(v))} \\
&= \sum_{v \in V_{m-1}} \frac{1}{2^{m-1}} \ln \frac{1}{2^{m-1} \Pr_{P_j}(v)} + \sum_{v \in V_{m-1}} \frac{1}{2^{m-1}} \left(\frac{1}{2} \ln \frac{1}{2 p_j(v, \ell(v))} + \frac{1}{2} \ln \frac{1}{2 p_j(v, r(v))} \right) \\
&= KL(D_F^{m-1}, D_{P_j}^{m-1}) + \sum_{v \in V_{m-1}} \frac{1}{2^{m-1}} KL((1/2, 1/2), (p_j(v, \ell(v)), p_j(v, r(v))))
\end{aligned}$$

Note that $KL((1/2, 1/2), (p_j(v, \ell(v)), p_j(v, r(v)))) = 0$ if $L(v) \neq j$ and otherwise it is at most $10\epsilon^2$ (earlier in the proof of Lemma 2 we show that $KL((1/2, 1/2), (1/2 + \epsilon, 1/2 - \epsilon)) \leq 10\epsilon^2$ for $\epsilon \leq 1/3$). So we have

$$KL(D_F^m, D_{P_j}^m) \leq KL(D_F^{m-1}, D_{P_j}^{m-1}) + \frac{1}{2^{m-1}} 10\epsilon^2 |v \in V_{m-1} : L(v) = j|$$

This implies that

$$\sum_{j=1}^K KL(D_F^m, D_{P_j}^m) \leq \sum_{j=1}^K KL(D_F^{m-1}, D_{P_j}^{m-1}) + 10\epsilon^2 \leq 10m\epsilon^2$$

where the last inequality in the above is by the repeated application of the first inequality. Now applying Pinsker inequality, we have

$$2 \sum_{j=1}^K d_{tv}^2(D_F^m, D_{P_j}^m) \leq 10m\epsilon^2 \implies \frac{\sum_{j=1}^K d_{tv}^2(D_F^m, D_{P_j}^m)}{K} < \frac{1}{200}$$

since we assumed that $m < \frac{K}{1000\epsilon^2}$.

Since the average of $d_{tv}^2(D_F^m, D_{P_1}^m), \dots, d_{tv}^2(D_F^m, D_{P_K}^m)$ is $1/200$ there exists at least $K/2$ values of $j \in [K]$ such that $d_{tv}^2(D_F^m, D_{P_j}^m) < 2 \cdot \frac{1}{200} = 1/100$. Let $i \neq k$ be two such values. We have $d_{tv}^2(D_F^m, D_{P_i}^m) < 1/100$ and $d_{tv}^2(D_F^m, D_{P_k}^m) < 1/100$. Hence $d_{tv}(D_F^m, D_{P_i}^m) < 1/10$ and $d_{tv}(D_F^m, D_{P_k}^m) < 1/10$. Since d_{tv} satisfies triangle inequality we have $d_{tv}(D_{P_k}^m, D_{P_i}^m) < 1/5$ which contradicts Claim 1.

Instance dependent lower bound

Recall that the expected regret of an algorithm depends on the instance. We will use $E[R(T, I)]$ for the expected regret of the algorithm on the instance I . Note that we have proved $\max_{I \in \mathcal{I}} E[R(T, I)] = O(\sqrt{KT \log T})$ for UCB/Successive Elimination where \mathcal{I} is the set of all bandit instances. Moreover, we have also seen the following instance-dependent upper bound for both UCB/Successive Elimination algorithms: $E[R(T, I)] = O(\log T) \sum_{a: \Delta_a > 0} \frac{1}{\Delta_a}$. We will now prove the matching instance dependent lower bound.

Theorem 13. *Let \mathcal{I} be the set of all bandit instances (where recall that a bandit instance consists of values of K, T and K reward distributions D_1, \dots, D_K).*

1. *There can not exist an algorithm for which $E[R(T, I)] \leq f(T)c_I$ for all bandit instance $I \in \mathcal{I}$ where $f(T) = o(\log T)$ and $c_I \geq 0$ depends on I only (not on T).*
2. *There can not exist an algorithm for which $E[R(T)] \leq \frac{\log T}{10000} \sum_{a: \Delta_a > 0} \frac{1}{\Delta_a}$ for all bandit instances.*

Proof. Let us first prove the first part. Suppose there exists such an algorithm B . Consider two bandit instances, both on 2 arms.

1. In the first instance, denoted by w , first arm is Bernoulli with a mean of $1/2$ and the second arm is a Bernoulli with a mean of 0.4 .
2. In the second instance, denoted by w' , first arm is Bernoulli with a mean of $1/2$ and the second arm is a Bernoulli with a mean of 0.6 .

Since all rewards are either 0 or 1, the run of the algorithm B on these two instances can be seen as a binary tree of height T in a natural way. All internal nodes of this tree are labelled by either 1 or 2 (representing the pull of the arm 1 or 2 respectively), while the leaves are unlabelled. From any internal node, the edge to left child is labelled by 0 and the edge to right child is labelled by 1 (this ordering is arbitrary). The algorithm starts from root and reach a leaf. When the algorithm reaches any internal node labelled by $i \in \{1, 2\}$ then it means that the algorithm pulls the arm i . If the reward received is 0 then the algorithm moves to the left child and if the reward received is 1 then it goes to the right child. In this way, algorithm starts from the root node and reaches some leaf node. Recall the notations $V_h, L(v), \dots$ etc. used for the tree from the previous proof as we will follow the same notations for clarity.

Let $E[R(T, w)]$ be the expected regret of the algorithm B on instance w after round T and $E[R(T, w')]$ be the expected regret of the algorithm B on instance w' after round T . For any $i \in \{1, 2\}$, let $E_w[N_i^T]$ is the expected number of times arm i is pulled by algorithm B after round T when the instance is w . Similarly we define $E_{w'}[N_1^T]$ and $E_{w'}[N_2^T]$.

Note that $E[R(T, w)] = 0.1 \cdot E_w[N_2^T]$ and $E[R(T, w')] = 0.1 \cdot E_{w'}[N_1^T]$. By assumption, we also have $E[R(T, w)] \leq f(T)c_w$ for all T and $E[R(T, w')] \leq f(T)c_{w'}$ for all T where $f(T) = o(\log T)$. Therefore, $E_w[N_2^T] \leq 10c_w f(T) =$

$o(\log T)$ (as c_w does not depend on T). Similarly, $E_{w'}[N_1^T] = o(\log T)$. Now we will use Markov inequality to get the following inequality:

$$\Pr_w(N_2^T \geq T/2) \leq \frac{2E_w[N_2^T]}{T} = o(\log T)/T$$

$$\Pr_{w'}(N_2^T \geq T/2) = 1 - \Pr_{w'}(N_1^T \geq T/2) \geq 1 - \frac{2E_{w'}[N_1^T]}{T} \geq 1 - o(\log T)/T$$

We can see that for instance w , the event $N_2^T \geq T/2$ happens with probability at most $o(\log T)/T$ (which tends to 0 as $T \rightarrow \infty$) whereas for the instance w' , the same event happens with probability at least $1 - o(\log T)/T$ (which tends to 1 as $T \rightarrow \infty$). This implies that the total variation distance between the distributions induced on leaves (recall algorithm B is just a tree) by instance w and w' respectively is large. Formally let D_w^T and $D_{w'}^T$ be the distribution induced on leaves for instance w and w' respectively. Let $S_1 \subseteq V_T$ be the set of leaves such that at least $T/2 + 1$ nodes in the path from the root to the leaf are labelled by 1 (in other words, if algorithm reaches any leaf in S_1 then we have $N_1^T > T/2$). Therefore $S_2 = V^T \setminus S_1$ is the set of leaves such that at most $T/2$ nodes in the path from the root to the leaf are labelled by 1 (in other words, if algorithm reaches any leaf in S_2 then we have $N_1^T \leq T/2$).

From Lemma 3,

$$d_{tv}(D_w^T, D_{w'}^T) \geq D_w^T(S_1) - D_{w'}^T(S_1) = 1 - \frac{o(\log T)}{T} - \frac{o(\log T)}{T} = 1 - \frac{o(\log T)}{T}$$

Recall that in all previous lower bounds proofs that we have seen, we have used Pinsker inequality $d_{tv}(P, Q) \leq \sqrt{\frac{1}{2}KL(P, Q)}$ to relate total variation distance and KL divergence. Unfortunately, this inequality is not useful when P and Q are very far from each other (because when P and Q are far, though tv distance is always at most 1, KL divergence between them can be a large number and hence above inequality will be vacuously true).

We will use the Bretagnolle–Huber bound: for any distributions P and Q ,

$$d_{tv}(P, Q) \leq \sqrt{1 - e^{-KL(P, Q)}} \leq 1 - \frac{1}{2}e^{-KL(P, Q)}$$

Note that $1 - \frac{1}{2}e^{-KL(P, Q)}$ is always < 1 and hence the above inequality is not vacuously true when $d_{tv}(P, Q) \rightarrow 1$ and $KL(P, Q) \rightarrow \infty$.

Now

$$KL(D_w^T, D_{w'}^T) \geq \ln \frac{1}{2(1 - d_{tv}(D_w^T, D_{w'}^T))} \geq \log \frac{T}{o(\log T)}$$

for all T . This implies that

$$\lim_{T \rightarrow \infty} KL(D_w^T, D_{w'}^T) / \log T \geq 1 \quad (2)$$

By the chain rule, we also have (as we did in previous lower bound proof)

$$KL(D_w^T, D_{w'}^T) = KL(D_w^{T-1}, D_{w'}^{T-1}) + \sum_{v \in V_{T-1}: L(v)=2} \Pr_w(v) KL((0.5, 0.4), (0.5, 0.6))$$

Let $d = KL((0.5, 0.4), (0.5, 0.6))$. By repeated application of the above inequality, we have

$$KL(D_w^T, D_{w'}^T) = d \sum_{h=0}^{T-1} \sum_{v \in V_h: L(v)=2} \Pr_w(v)$$

Note that for any h , $\sum_{v \in V_h: L(v)=2} \Pr_w(v)$ is the probability that the algorithm B pulls arm 2 in the round $h+1$ (for instance w). Thus $\sum_{h=0}^{T-1} \sum_{v \in V_h: L(v)=2} \Pr_w(v)$ is just $E_w[N_2^T]$ (hint: think of indicator random variable). Now

$$\begin{aligned} \lim_{T \rightarrow \infty} KL(D_w^T, D_{w'}^T) / \log T &\geq 1 \quad (\text{from inequality 2}) \\ \implies \lim_{T \rightarrow \infty} \frac{E_w[N_2^T]}{\log T} &\geq \frac{1}{d} \end{aligned}$$

Recall that earlier we have shown that $E_w[N_2^T] = o(\log T)$ which is a contradiction to above. \square

The proof of second part is almost same. Modify instances w and w' as follows and rest of the proofs follows in similar fashion:

1. In w , first arm is Bernoulli with a mean of $1/2$ and the second arm is a Bernoulli with a mean of $0.5 - \epsilon$.
2. In w' , first arm is Bernoulli with a mean of $1/2$ and the second arm is a Bernoulli with a mean of $0.5 + \epsilon$.

6 Adversarial bandits

In adversarial bandits, there is no randomness in the received rewards. The input is an unknown but fixed reward table (which is just a sequence of T reward vectors $\mathbf{r}_1, \dots, \mathbf{r}_T$ where $\mathbf{r}_i = (r_{i,1}, r_{i,2}, \dots, r_{i,K}) \in [0, 1]^K$). As there is no randomness in the received rewards, the randomness can come from the algorithm only.

Input/Instance: K, T (known), unknown $\mathbf{r}_1, \dots, \mathbf{r}_T$
In each round $t \in T$:

1. algorithm picks an arm $a_t \in [K]$.
2. algorithm receives the reward r_{ta_t} .

How do we define regret? One way to define regret is as follows : regret of an algorithm is the difference of the total rewards received by the algorithm that knows the reward table minus the (expected) reward received by the algorithm. The issue with this definition is that it makes the problem very hard in the sense that any algorithm will have large regret ($\Omega(T)$).

The regret of an algorithm is defined in the following way:

$$R(T) = \max_a \sum_{t=1}^T r_{ta} - \sum_{t=1}^T r_{ta_t}$$

That is, the regret of an algorithm is the total rewards received by the best arm in hindsight minus the total reward received by the algorithm. There are several advantages of defining the regret in the above mentioned way:

1. Now we will be able to show algorithms with sublinear regret.
2. This can be starting point to study more stronger notions of regret (something in between the above two notions).

As it will be clear soon, there is good reason to consider costs instead of rewards with the goal of minimizing cost (than maximizing rewards). Note that it is just a matter of convenience and to be in line with existing literature we will consider costs.

Input/Instance: K, T (known), unknown cost vectors $\mathbf{c}_1, \dots, \mathbf{c}_T$ where $\mathbf{c}_i = (c_{i1}, \dots, c_{iK})$
In each round $t \in T$:

1. algorithm picks an arm $a_t \in [K]$.
2. algorithm pays the cost c_{ta_t} .

The regret of an algorithm is defined in the following way:

$$R(T) = \sum_{t=1}^T c_{ta_t} - \min_a \sum_{t=1}^T c_{ta}$$

Bandits with full feedback

An easier problem is bandits with full feedback where after every round t , the entire cost vector \mathbf{c}_t is revealed. Note that this is easier than adversarial bandits. As we will see later, this is also a special case of online learning.

Input/Instance: K, T (known), unknown cost vectors $\mathbf{c}_1, \dots, \mathbf{c}_T$
In each round $t \in T$:

1. algorithm picks an arm $a_t \in [K]$ and pays the cost c_{ta_t} .
2. algorithm gets to know the cost vector \mathbf{c}_t .

A general idea is to first study/give algorithms for bandits with full feedback. Then modify the algorithm to make it work for bandit setting.

Following theorem shows that any deterministic algorithms will have large regret even for bandits with full feedback.

Theorem 14. *Any deterministic algorithm will have a regret of $T/2$ for some instance, even for 2 arms and when costs are either 0 or 1.*

Proof Sketch: Fix any deterministic algorithm. Construct an instance consisting of only 2 arms and binary cost 0 or 1 as follows: in any round t , exactly one arm has cost of 1 and other has cost of 0 (that is, for any t , either $c_{t1} = 1, c_{t2} = 0$ or $c_{t1} = 0, c_{t2} = 1$). Depending on the algorithm, the cost table is constructed so that algorithm pays the total cost of T (fill the details yourself) while there exists an arm that pays only a cost of at most $T/2$ (this is directly implied by the fact that in any round, one arm pays cost of 0 and other 1). \square

Note that this is in sharp contrast with stochastic bandits, where all the algorithms that we had seen were deterministic. Now we give a randomized algorithm Hedge that has sublinear regret for bandits with full feedback.

Hedge Algorithm

Theorem 15. *For any $\epsilon > 0$ and cost vectors $\mathbf{c}_1, \dots, \mathbf{c}_T$ with non negative entries, the expected regret of Hedge algorithm satisfies:*

$$E[R(T)] \leq \frac{\log K}{\epsilon} + \epsilon \sum_{t=1}^T \mathbf{p}_t \cdot \mathbf{c}_t^2$$

where $\mathbf{c}_t^2 = (c_{t1}^2, \dots, c_{tK}^2)$ and $\mathbf{p}_t \cdot \mathbf{c}_t^2 = \sum_a p_t(a) c_{ta}^2$.

Corollary 2. *If each cost is in $[0, 1]$ (that is, $c_{ta} \in [0, 1]$ for each t and a) then the expected regret of Hedge satisfies:*

$$E[R(T)] = O(\sqrt{T \log K})$$

Algorithm 6: Hedge Algorithm

```
Parameter :  $\epsilon > 0$  (will be chosen to optimize regret);  
 $w_1(a) = 1$  for all arms  $a$ ;  
/* Initialization */  
for  $t = 1, \dots, T$  do  
     $p_t(a) = w_t(a) / \sum_a w_t(a)$ ;  
     $\mathbf{p}_t = (p_t(1), \dots, p_t(K))$ ;  
     $a_t \sim \mathbf{p}_t$ ;  
    Algorithm pays the cost  $c_{ta_t}$ ;  
    Algorithm receives the cost vector  $\mathbf{c}_t$ ;  
     $w_{t+1}(a) = w_t(a)e^{-\epsilon c_{ta}}$  for all  $a$  ;  
    /* update of weights */  
end
```

Proof. If each cost is in $[0, 1]$ then we have $\mathbf{p}_t \cdot \mathbf{c}_t^2 \leq 1$ for each t . Hence $E[R(T)] \leq \frac{\log K}{\epsilon} + \epsilon T$. Choosing $\epsilon = \sqrt{\frac{\log K}{T}}$ gives $E[R(T)] = \sqrt{T \log K}$. \square

For the proof of theorem, follow this material (beware of change in notations):

https://people.eecs.berkeley.edu/~jiantao/2902021spring/scribe/EE290_Lecture_10.pdf

EXP3 Algorithm

The idea of EXP3 algorithm is to simulate the Hedge algorithm on fake cost vectors. In any round t , the algorithm pulls arm a_t and hence gets to know only c_{ta_t} . We define fake cost vector $\hat{\mathbf{c}}_t = (0, \dots, \frac{c_{ta_t}}{p_t(a_t)}, 0, \dots, 0)$ (all entries are 0 except the a_t th one which is $\frac{c_{ta_t}}{p_t(a_t)}$). EXP3 algorithm just invokes Hedge on the fake cost vector to update weights and proceed further.

Algorithm 7: EXP3 Algorithm

Parameter : $\epsilon > 0$ (will be chosen to optimize regret);
 $w_1(a) = 1$ for all arms a ;
/* Initialization */
for $t = 1, \dots, T$ **do**
 $p_t(a) = w_t(a) / \sum_a w_t(a)$;
 $\mathbf{p}_t = (p_t(1), \dots, p_t(K))$;
 $a_t \sim \mathbf{p}_t$;
 Algorithm pays the cost c_{ta_t} ;
 /* fake cost vector $\hat{\mathbf{c}}_t = (0, \dots, \frac{c_{ta_t}}{p_t(a_t)}, 0, \dots, 0)$ */
 For each arm a : if $a \neq a_t$ then $w_{t+1}(a) = w_t(a)$ else if $a = a_t$ then
 $w_{t+1}(a) = w_t(a) e^{-\epsilon \frac{c_{ta_t}}{p_t(a_t)}}$;
 /* update of weights */
end
